

Coursework: Analysing deep neural networks as brain models

For this coursework you are asked to implement and explore the behaviour of backprop using supervised learning and a deep neural network as a potential model for information processing in the brain. For implementation we suggest that you use lab1 as a starting point and Python or Julia in particular.

1. Implement the backpropagation algorithm in the supervised context

The backpropagation algorithm is often used in machine learning to solve the credit assignment problem. Here you are going to contrast its different elements such as error backpropagation, or symmetric weights, on how the brain is organised. You should use a simple feedforward neural network with one or two hidden layers and sigmoidal units to teach the network to classify the widely used fashion MNIST dataset (<https://github.com/zalandoresearch/fashion-mnist>); to simplify the problem you can reduce the dataset to five of the ten labels. Do not use autograd tools as in pytorch for this simple task. You should use a mean squared cost as discussed in the lectures and a one-hot output¹ encoding to represent the fashion MNIST classes. Include brief snippets of your code in the report. Note: Although here you are asked to implement the backprop algorithm in a supervised setting, backprop is also used in unsupervised (e.g. autoencoders) and reinforcement learning (e.g. deep RL).

You should use the lectures and you may also use these or related papers as references when contrasting backprop with the brain:

- Blake and Lillicrap, Dendritic solutions to the credit assignment problem, Current Opinion in Neurobiology 2019.
- Sacramento et al., NeuralIPS 2018, ([arXiv.org:1810.11393](https://arxiv.org/abs/1810.11393))

2. Analyse the deep network using information theory

Recently information theory has been used to study the dynamics of learning in deep learning networks; here you will examine the mutual information

¹<https://en.wikipedia.org/wiki/One-hot>

between the image label and the network activity: you can imagine a communication channel

label \rightarrow image \rightarrow hidden layer \rightarrow output

in which the image label is encoded in the image itself and that in turn is encoded in the activity of the hidden layer and then in the output layer. The aim is to quantify information flow in this channel, at least in a discrete approximation in which neurons are either on or off. Thus a random variable representing a layer will take binary-sequence values with ones and zeros corresponding to neurons being “on” or “off” based on their activation. The probability for each binary sequence can be estimated by counting the number of images that produce that sequence; note that for the hidden layer many will have zero estimated probability.

- Opening the black box of deep neural networks via information. Ravid Shwartz-Ziv and Naftali Tishby. [arXiv:1703.00810](#)
- Adaptive estimators show information compression in deep neural networks. Ivan Chelombiev, Conor Houghton and Cian O'Donnell. [arXiv:1902.09037](#)

Questions

1. Biological relevance of backprop
 - (a) *How does the algorithm work?* Explain the algorithm (with equations) and plot the different components over learning to help you explain its behaviour, for example what are the weight changes given by backprop and the different terms that make up the gradient. You should also plot the performance (i.e. cost) of the algorithm over learning iterations (i.e. the learning curve). You should include a snippet of the key component of your code in your report. [3/20 marks, max 500 words]
 - (b) *How does the algorithm relate to the brain?* Discuss how backprop relates to the brain. Can neural networks optimised with backprop explain neuroscience data? Here you are going to explore the biological plausibility of three of the key issues with backprop (you should discuss their implications and how important they

are using simulations and the respective plots to support your arguments) [4/20 marks, max 500 words]:

- i. Weight transport problem
 - ii. The need for derivative of the activation function
 - iii. Two phase learning
- (c) Discuss what are the key advantages of using supervised learning over the other two learning paradigms, unsupervised and reinforcement learning, both in terms of performance/data needs and biological plausibility. For this [1.5/20 marks, max 250 words]
- (d) What about disadvantages in terms of performance/data needs and biological plausibility? And how could this be improved upon? [1.5/20 marks, max 250 words]

2. Information theory analysis

- (a) If X is a random variable representing the input, $H(X)$ will be $\log(5)$ since there are five labels or $\log(10)$ if you are using the whole data set. Let Y represent the discretized hidden layer activity, this means the components of Y represent the activity of individual nodes of the hidden layer with an integer. One approach might be to find the average activity and use that as a cut-off with values exceeding it mapped to one and other values to zero. Now, what is $H(Y|X)$ and how does this change with learning? What about $I(X, Y)$? How are these dynamics changed if the number of hidden layer neurons is changed? If Z represents the activity of the output layer, how do $H(Z|X)$ and $I(X, Z)$ change? [5/20 marks, maximum 500 words and four graphs]
- (b) What happens if you consider some property of the image, for example, if the image is divided in four and W is the random variable giving which quadrant has the largest amount of white. How do $I(W, Y)$ and $I(W, Z)$ change with learning? [3/20 marks, maximum 250 words and two graphs]
- (c) As a bonus question, you could consider the performance of an autoencoder. How does the mutual information of the image properties change with learning in this case? [2/20, maximum 250 word and two graphs]

Submit your report on Blackboard (IPB > Assessment, submission and feedback) remembering that you need to submit to the ‘coursework’ unit (COMSM0100) not the ‘teaching’ unit (COMSM0075). No need to submit your full code (include snippets in the report), only the report in pdf or similar. Please add the necessary references/citations at the end.

Note 1 : Where possible cite papers and/or use simulations/plots to support your claims.

Note 2 : Collaborative work is encouraged for example, for coding and understanding of the algorithm, but every submission should be individual / unique.

Support provided

We cannot provide direct support on the coursework, but we (TAs and Lecturers) can help with questions you might have about the lectures and labs/workshops. We will be available during the usual lab/workshop times as in your timetable though we may sometimes only be available online.

Deadline

The deadline for submission of all optional unit assignments is 13:00 on Thursday 8th of December: the University discourages Friday deadlines. Students should submit all required materials to the “Assessment, submission and feedback” section of Blackboard - it is essential that this is done on the Blackboard page related to the “With Coursework” variant of the unit.

Time commitment and marking

You are expected to work on both of your optional unit courseworks in the three-week coursework period as if it were a working week in a regular job - that is five days a week for no more than eight hours a day. The effort spent on the assignment for each unit should be approximately equal, being roughly equivalent to 1.5 working weeks each. It is up to you how you distribute your time and workload between the two units within those constraints.

You are strongly advised NOT to try and work excessive hours during the coursework period: this is more likely to make your health worse than to make your marks better. If you need further pastoral/mental health support,

`comsm0075.github.io`

please talk to your personal tutor, a senior tutor, or the university wellbeing service.

Academic Offences

Academic offences, including submission of work that is not your own, falsification of data/evidence or the use of materials without appropriate referencing, are all taken very seriously by the University. Suspected offences will be dealt with in accordance with the University's policies and procedures. If an academic offence is suspected in your work, you will be asked to attend an interview with senior members of the school, where you will be given the opportunity to defend your work. The plagiarism panel is able to apply a range of penalties, depending on the severity of the offence. These include: requirement to resubmit work, capping of grades and the award of no mark for an element of assessment.

Extenuating circumstances

If the completion of your assignment has been significantly disrupted by serious health conditions, personal problems, periods of quarantine, or other similar issues, you may be able to apply for consideration of extenuating circumstances (in accordance with the normal university policy and processes). Students should apply for consideration of extenuating circumstances as soon as possible when the problem occurs, using the following online form: <https://www.bristol.ac.uk/request-extenuating-circumstances-form>

You should note however that extensions of any significant length are not possible for optional unit assignments. If your application for extenuating circumstances is successful, you may be required to retake the assessment of the unit at the next available opportunity (for example, during the summer reassessment period).